

pse-edge-mcp — Developer Walkthrough

Audience: a developer or architect who has never seen this codebase and needs to understand it, debug it, or extend it.

Companion documents. This file explains *how the system works*. Two others explain *why it is the way it is*, and you should read them before changing anything structural:

Document	Contents
docs/plan.md	Every design decision: scope, caching policy, auth design, phases, risks
docs/endpoints.md	The verified PSE Edge endpoint map — request dialects, parameter names, response shapes
docs/deploy.md	Running it in production
CLAUDE.md	The short form of the invariants, kept next to the code

Version described: **0.9.0**. 35 modules, ~7,500 lines of source, ~5,600 lines of tests.

1. What this is, in one paragraph

An **MCP server** (Model Context Protocol — a standard way to expose tools to an LLM client) that serves **Philippine Stock Exchange** data taken from the PSE Edge disclosure portal. It exposes **12 read-only tools** — quotes, price history, disclosures, financial reports, dividends, index and market data — plus **one action tool**, `send_email`, on deployments with auth enabled.

It is **unofficial**. PSE Edge publishes no API, so this server speaks to the same internal endpoints the portal's own web pages call.

The one decision that explains the whole design

PSE Edge is a public utility with no rate-limit contract and no commercial relationship with this project. Hammering it during trading hours would be antisocial and would get the project blocked. So:

This is an end-of-day server. It makes zero upstream requests while the market is open.

Almost every structural choice follows from that sentence — the cache is a *freeze* rather than a TTL, every response carries freshness metadata, there is an opportunistic archive, and there is a politeness throttle independent of user quotas. If a change you are contemplating would increase load on PSE Edge, it is probably the wrong change.

2. Running it

Three modes, same tool implementations underneath.

MODE	COMMAND	AUTH	STORAGE
studio (local)	uvx pse-edge-mcp (Claude Desktop / Claude Code default)	never	in-memory
HTTP (dev)	pse-edge-mcp --http --port 8000	opt-in	in-memory or Postgres
HTTP (production)	docker compose -f compose.nas.yaml up -d + compose.tunnel.yaml	on	Postgres

studio never authenticates. It runs on the user's own machine as a subprocess of their client; there is no network boundary to guard and no second party to authenticate. Adding auth there would be pure friction.

Local development:

```
uv sync --all-extras      # Python 3.14, uv for everything
uv run pytest              # 287 tests, no network access
uv run ruff check .        # line length 100
uv run mypy src            # strict
```

3. Request lifecycle

This is the path every tool call takes. Read it once and most of the codebase becomes predictable.

MCP client (Claude Desktop, claude.ai connector, your own agent)

POST /mcp {"jsonrpc":"2.0","method":"tools/call", ...}



HealthApp	/health, /health/ready – answered here, never authenticated, never touch the DB
AuthApp	/oauth/*, /signup, /login, /account, /privacy, / – reachable WITHOUT a token, because you cannot present one before you have one
AuthMiddleware	bearer token → user; quota check; 401 + WWW-Authenticate, or 429
MCP app (SDK)	JSON-RPC framing, tool dispatch, DNS-rebinding Host check



server.py validate args → call repository → shape reply
 (no domain logic here – see §5)



repositories.py build cache key → FreezeService.get(...) → parse → model



service.py
FreezeService THE FREEZE DECISION – see §4
 cache hit? market open? fetch or refuse



 cache hit cache miss + closed

Storage client.py —HTTP→ edge.pse.com.ph
(memory | Postgres) parsers.py ← HTML / JSON



 ↘ archive.py (opportunistic, best-effort)

 { "data": ..., "meta": {as_of, valid_until, from_cache, stale, data_policy}}

Two things worth noticing:

- **The composition is built exactly once**, in `asgi.py`. The CLI (`__main__.py`) and production both call `create_app()`, so they cannot drift apart.
- **/mcp is the only authenticated path**. Everything the browser needs during signup is deliberately in front of the auth gate.

4. The freeze policy — the core behaviour

This is the part to understand before anything else, because it is what makes the server *correct* rather than merely functional.

`FreezeService.get()` in `service.py` implements one decision table:

	MARKET CLOSED	MARKET OPEN
cache fresh	serve from cache	serve from cache
cache stale (expired)	FETCH upstream, store, serve	serve the stale value, flagged <code>meta.stale = true</code> ← never fetches
cache empty	FETCH upstream, store, serve	raise <code>MARKET_OPEN_NO_CACHE</code> with a <code>retry_after</code> timestamp

Market hours are **09:30–15:00 Asia/Manila on trading days**; weekends and the `PSE_HOLIDAYS` table in `market_calendar.py` are non-trading.

Consequences a newcomer will trip over

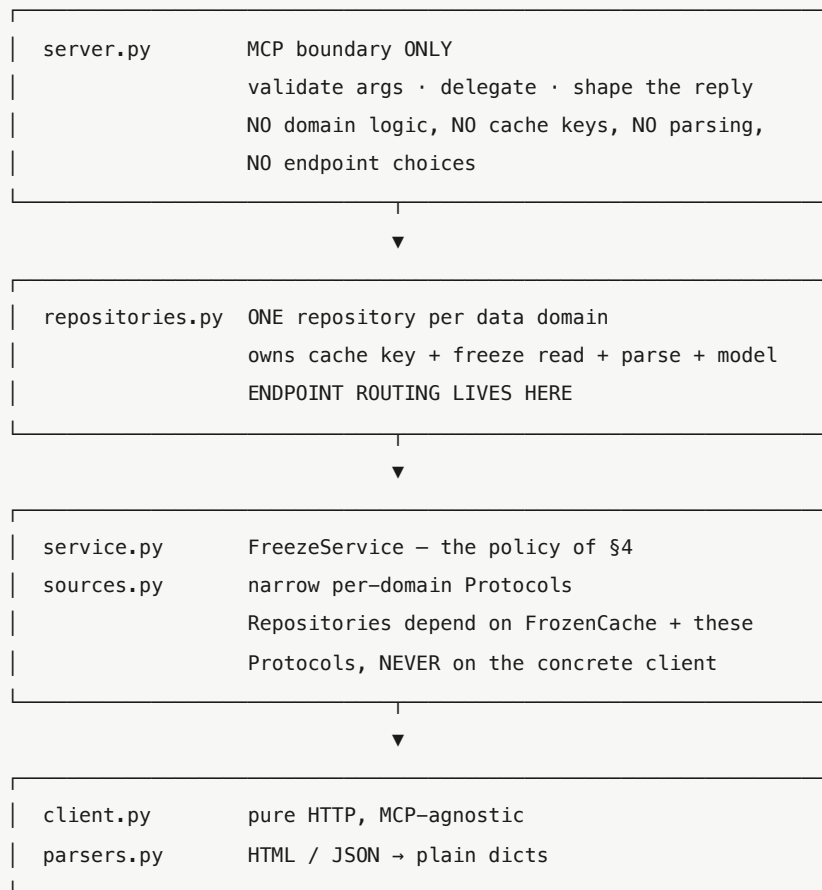
- **`MARKET_OPEN_NO_CACHE` is not a bug.** It is the policy working. Agents must handle it; a `retry_after` timestamp is included so they can schedule rather than poll.
- **`meta.stale = true` means the market is open** and you are seeing the last close. It does not mean the data is wrong.
- **Concurrent cache misses collapse into one upstream request** (`SingleFlight` in `ratelimit.py`). Ten simultaneous first-time requests for the same symbol produce one HTTP call, not ten.
- **`immutable=True`** (passed by the disclosure-detail path) marks objects that never change upstream — a disclosure identified by `edge_no`. Those are cached forever with `valid_until: null` and `data_policy: "immutable"`. The open-market freeze still governs their *first* fetch: protecting PSE Edge outranks serving a cache miss promptly.

The politeness throttle is separate from quotas

`throttle_rate_per_sec` (default 1.0/s, burst 2) limits **outbound** requests to PSE Edge and exists to protect *them*. Per-user quotas (60/min, 2000/day) limit **inbound** requests and exist to protect *you*. Do not conflate them; raising one does not justify raising the other.

5. Architecture and layering

The layering is enforced by convention and by tests. Violating it is the most likely way a well-meaning change gets rejected in review.



Supporting modules: `validation.py` (argument validators), `models.py` (Pydantic), `errors.py`, `cache.py` (Storage protocol), `archive.py`, `ratelimit.py`, `market_calendar.py`, `config.py`, `memo.py`.

Why repositories depend on Protocols

`sources.py` defines five narrow interfaces — `CompanySource`, `QuoteSource`, `DisclosureSource`, `CompanyInfoSource`, `MarketSource`. `PseEdgeClient` happens to satisfy all of them, but repositories only ever see the narrow one they need. The payoff is in `tests/test_repositories.py`: a repository can be tested with a five-line fake and **no HTTP mocking at all**.

Module map

Module	Lines	Responsibility
<code>auth_app.py</code>	756	Browser- and client-facing routes: OAuth, signup, login, account, privacy
<code>parsers.py</code>	752	HTML/JSON → dicts. The most PSE-specific code in the repo
<code>oauth.py</code>	721	OAuth 2.1 server: DCR, PKCE, code exchange, refresh rotation, <code>client_credentials</code>
<code>repositories.py</code>	526	The domain layer — five repositories
<code>admin.py</code>	431	<code>pse-edge-admin</code> CLI, including machine-client provisioning
<code>passkeys.py</code>	424	WebAuthn ceremonies and browser sessions
<code>server.py</code>	413	The 12 tool definitions
<code>models.py</code>	332	24 Pydantic models
<code>client.py</code>	268	HTTP to PSE Edge; both request dialects
<code>db.py</code>	255	SQLAlchemy tables, engine, schema check
<code>asgi.py</code>	240	The single composition point, and the DNS-rebinding allowlist
<code>auth.py</code>	190	<code>TokenService</code> , <code>QuotaTracker</code> . Must stay SQLAlchemy-free
<code>notifications.py</code>	143	<code>send_email</code> policy: self-only recipient, caps, escaping

6. The tool surface

Twelve tools are read-only and return the same envelope (§7). The thirteenth, `send_email` , is an **action** — see below.

Tool	Arguments	Repository
<code>search_companies</code>	<code>query</code>	<code>CompanyRepository.search</code>
<code>validate_symbol</code>	<code>symbol</code>	<code>CompanyRepository.try_resolve</code>
<code>get_stock_quote</code>	<code>symbol</code>	<code>QuoteRepository.quote</code>
<code>get_price_history</code>	<code>symbol</code> , <code>start_date?</code> , <code>end_date?</code>	<code>QuoteRepository.history</code>
<code>search_disclosures</code>	<code>symbol?</code> , <code>start_date?</code> , <code>end_date?</code> , <code>template?</code> , <code>page</code>	<code>DisclosureRepository.search</code>
<code>search_disclosure_fulltext</code>	<code>keyword</code> , <code>start_date?</code> , <code>end_date?</code> , <code>symbol?</code> , <code>subject_title?</code> , <code>page</code>	<code>DisclosureRepository.fulltext</code>
<code>get_disclosure</code>	<code>edge_no</code>	<code>DisclosureRepository.detail</code>
<code>get_company_profile</code>	<code>symbol</code>	<code>CompanyInfoRepository.profile</code>
<code>get_financial_highlights</code>	<code>symbol</code>	<code>CompanyInfoRepository.financials</code>
<code>get_dividends_and_rights</code>	<code>symbol</code>	<code>CompanyInfoRepository.dividends_and_rights</code>
<code>get_indices</code>	—	<code>MarketRepository.indices</code>
<code>get_market_summary</code>	—	<code>MarketRepository.summary</code>
<code>send_email</code> (<i>auth only</i>)	<code>subject</code> , <code>body</code>	<code>NotificationService.send_to_self</code>

Dates are `YYYY-MM-DD` on the tool surface; the client converts to PSE Edge's `MM-dd-yyyy` wire format.

`search_disclosures` routes to two different upstream endpoints depending on whether `symbol` is supplied — market-wide with a date range, or one company's full history. That routing decision lives in the repository, which is exactly where the layering says it belongs.

`search_disclosure_fulltext` is deliberately a separate tool, not a better search. PSE Edge's keyword index only covers roughly 2023–2025, so it is not a superset of `search_disclosures`. The tool reports this coverage limit in its own results.

`send_email` — the one action tool

Policy lives in `notifications.py`, away from the MCP boundary, so every rule is testable without HTTP or a mail provider. Read that module's docstring before changing anything here.

The recipient is not a parameter. It is resolved from the ASGI scope the auth middleware populated (`server._caller`), i.e. from a validated bearer token. That single decision is what makes a mail tool safe to expose on a public server:

- **No open relay.** Signup is self-service, so a `to` argument would let anyone on the internet send mail from the operator's domain. The domain gets blocklisted, the provider suspends the account for abuse — and since the same provider sends verification email, *signup breaks with it*.
- **Nothing for prompt injection to steer.** This server returns disclosure text fetched from PSE Edge — third-party content the model reads and nobody here controls. "Email this to `attacker@example.com`" planted in a disclosure is an exfiltration path for any tool that accepts an address. Here there is no address argument to poison.

The rest is blast radius, since the direction is already fixed: registered **only when auth is enabled** (an auth-less deployment has no verified address, and a tool that is always listed and always fails just makes a model keep choosing

it); body **escaped, never rendered as HTML**, because model-authored markup arriving from a trusted domain is what phishing looks like; 200-character subject, 20,000-character body, 20 messages per user per day.

It returns `{"data": ...}` with **no meta**, via `act()` rather than `reply()`. `meta` answers "how fresh is this market data"; an action has no `as_of` and no `valid_until`, and inventing one would quietly make the freshness contract meaningless where it does matter.

7. The data contract

Every tool returns this envelope

```
{
  "data": { "...": "the payload" },
  "meta": {
    "as_of":      "2026-07-31T15:00:00+08:00",
    "valid_until": "2026-08-01T15:00:00+08:00",
    "from_cache": true,
    "stale":      false,
    "data_policy": "EOD-frozen"
  }
}
```

`meta` is not decoration. It is the mechanism by which the server is honest about a freeze policy that would otherwise look like stale data. A client that ignores `meta` will eventually mislead someone about when a price was true.

- `data_policy: "immutable"` with `valid_until: null` marks objects that never change upstream.
- **Clients should cache against `valid_until`**. Re-querying before it passes returns a byte-identical answer and only burns quota.

Error codes

Errors are returned as data — `{"error": "CODE", "message": ..., ...}` — not as exceptions, so an LLM can react to them rather than seeing a traceback.

Code	Meaning	Caller should
MARKET_OPEN_NO_CACHE	Nothing cached and the market is open	Retry after <code>retry_after</code>
SYMBOL_NOT_FOUND	No such ticker	Try <code>search_companies</code>
INVALID_ARGUMENT	Failed validation	Fix the arguments
ENDPOINT_CHANGED	PSE Edge's response shape changed	Alert a human — see §11
EDGE_UNAVAILABLE	Upstream unreachable or erroring	Retry later
RATE_LIMITED	An action's own budget is spent (e.g. daily email cap)	Honour <code>retry_after_seconds</code>
ACTION_UNAVAILABLE	The action cannot run here (no authenticated caller, or no mail provider)	Stop asking — do not retry
INTERNAL_ERROR	Anything else	File a bug

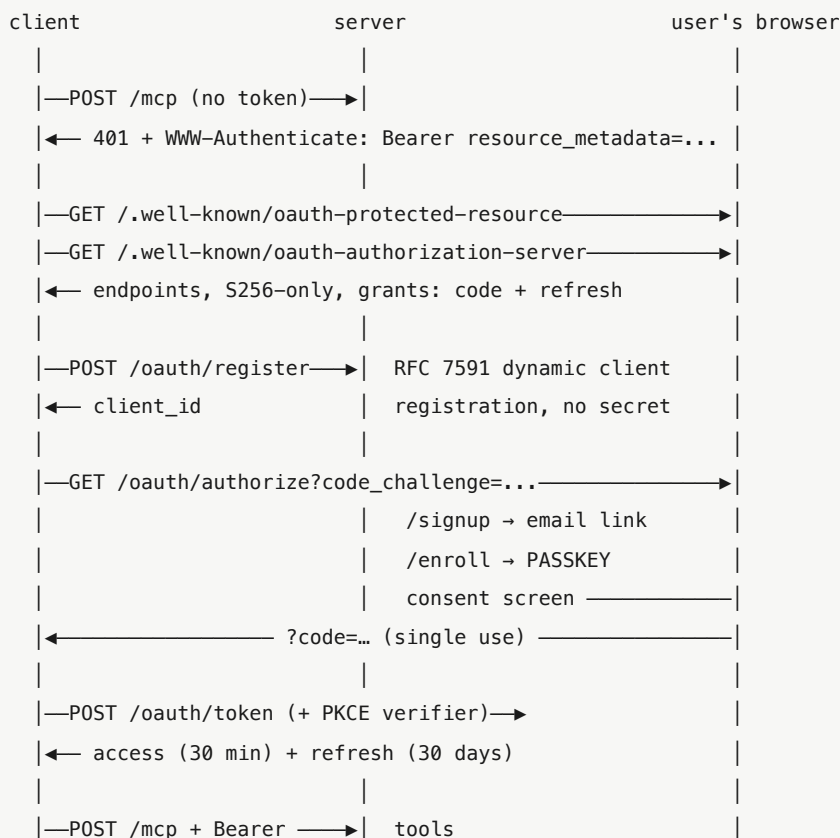
The mapping happens once, in `server.reply()`. Do **not** add per-tool `try/except`.

`reply()` is a helper that takes a callable, and deliberately not a decorator. The MCP SDK builds each tool's output schema from its return annotation, and a `functools.wraps` wrapper makes `inspect.signature` follow `__wrapped__` to the *inner* annotation, breaking schema generation. Tools must stay annotated `-> dict[str, Any]`. `tests/test_server_disclosures.py::test_tool_surface_is_stable` guards this.

8. Authentication

Auth is **opt-in** (`PSE_AUTH_REQUIRED=1`), which requires `DATABASE_URL`). `stdio` never authenticates.

The self-service flow



Exactly two steps involve a human: the passkey enrollment and the consent click. No passwords exist anywhere in the system.

Rules that must not regress

- Unknown client or unregistered redirect URI is **fatal and never redirects** — otherwise the server is an open redirector.
- Redirect URIs match **exactly**. No prefix matching.
- PKCE is mandatory, S256 only.**
- Authorization codes are single-use via one atomic `UPDATE ... WHERE consumed_at IS NULL RETURNING`.
- Refresh tokens rotate; **replaying a rotated one revokes the entire family** (RFC 9700 §4.14). A stolen refresh token gets one use before the theft is detected.
- Only `kind='access'` authenticates as a bearer.
- Tokens are opaque, `pse_`-prefixed, stored **only as SHA-256**.

- The validation-cache TTL (default 60 s) is **the revocation-latency budget and nothing else**. Refusals are never cached.
- `auth.py` must stay **SQLAlchemy-free** so the lean install path works; Postgres code lives in `auth_store.py` and imports lazily.

Headless and agentic access — `client_credentials`

A LangGraph app or the Anthropic Messages API MCP connector cannot open a browser, so it exchanges a client id and secret for a bearer token directly:

```

agent —POST /oauth/token—> grant_type=client_credentials
    |
    | HTTP Basic, or client_id/client_secret in the body
    |
    ← access_token (1 h), token_type=Bearer, scope=mcp
    |
    | ↑ NO refresh token
    |
    —POST /mcp + Bearer —> tools

```

Provisioned out of band; this is deliberately not reachable over HTTP:

```

pse-edge-admin create-machine-client --name langgraph-app # secret shown ONCE
pse-edge-admin revoke-machine-client <client_id>

```

The security property to understand before touching any of this: `/oauth/register` is open to the internet, so the right to use `client_credentials` must not be derivable from anything a registrant supplies — not a `grant_types` array it declares, not a requested auth method, not the presence of a secret. It is gated on `oauth_clients.client_type == 'machine'`, a column **only the admin CLI writes**. A DCR client that registers, declares the grant and sends a secret still gets `unauthorized_client`. If that check ever becomes conditional on request data, anyone on the internet can mint tokens.

Two more decisions worth knowing before changing them:

- **A machine client is backed by a service user** (under the reserved, non-routable `machine.invalid` domain). That is what keeps the bearer path identical for both grants: `/mcp` validates by joining `auth_tokens` to `users`, so a token with no user behind it would need a special case in the middleware — and a second path through authentication is where a revocation check goes missing later. It also gives machine clients quotas, usage accounting and disablement for free.
- **No refresh token for this grant.** The client already holds a long-lived secret it can present again; a refresh token would be a second credential of equal power without the rotation benefit that justifies one.

`POST /oauth/token` is rate-limited (`FixedWindowLimiter` in `ratelimit.py`, 20/minute per `client_id` **and** per IP), because it is the one endpoint where a long-lived credential can be attacked online. Both keys are always counted, so tripping one cannot keep the other cold.

The older route still exists for clients that speak neither OAuth nor this grant, and is the only option on a plain-HTTP deployment since browsers restrict WebAuthn to secure contexts:

```

pse-edge-admin create-user agent@example.com
pse-edge-admin issue-token agent@example.com --note nightly-job # shown once

```

Either way, give each agent **its own client or user**: quotas are per user, so a runaway job throttles itself, and revocation is surgical.

Privacy surface

`/account` shows a signed-in user everything held about them; `POST /account/delete` erases it immediately, in one transaction, with no approval step. Usage is aggregated **per user-hour, never per request** — minimal collection *and* no write on the hot path — and purged after 90 days.

`tests/test_privacy.py::test_erasure_leaves_nothing_behind` walks `metadata.tables`, so a new user-keyed table **fails that test** rather than silently retaining personal data. If you add a table with a user foreign key, that test is your reminder to wire it into erasure.

9. Storage and the archive

Postgres is **optional**. `DATABASE_URL` unset gives an in-memory cache and a `NullArchive` — the zero-config stdio path, which does not even need the `postgres` extra installed.

<code>DATABASE_URL</code> unset	<code>DATABASE_URL</code> set
<code>InMemoryStorage</code>	<code>PostgresStorage</code> → shared cache, so
<code>NullArchive</code>	<code>N</code> replicas still cause exactly one
(per-process, ephemeral)	upstream fetch per boundary
	<code>PostgresArchive</code> → EOD bars and
	disclosures accumulate over time

Tables: `cache_entries`, `eod_bars`, `disclosures`, `users`, `auth_tokens`, `webauthn_credentials`, `web_sessions`, `oauth_clients`, `oauth_flows`, `email_verifications`, `usage_events`.

Two columns carry more weight than their size suggests. `oauth_clients.client_type` (`'dcr'` | `'machine'`) is the entire authorization boundary for `client_credentials` — see §8. `oauth_clients.service_user_id` links a machine client to the account its tokens authenticate as, deliberately without a foreign key, matching the `oauth_flows.user_id` precedent so erasure ordering stays simple.

The archive is opportunistic. Nothing crawls. It fills only from fetches a user already caused, so it deepens over time at zero additional cost to PSE Edge — which matters because PSE Edge itself serves only limited history.

Rules:

- **Schema is applied by Alembic only**, never `create_all` at runtime. Compose runs a one-shot `migrate` service and `check_schema()` fails loudly at startup if migrations were skipped.
- **A failed archive write must never fail a read.** It is caught broadly — a dead database raises `OSError`, not merely `SQLAlchemyError`.
- Postgres imports are **lazy**, inside `build_storage()`. A test asserts SQLAlchemy does not leak into `sys.modules` on `import pse_edge_mcp.server`.

10. Extending the server

Adding a tool to an existing domain

1. Add a method to the relevant repository in `repositories.py`. It owns the cache key, the `FreezeService.get()` call, the parse, and the model.

2. Add a thin tool in `server.py` : validate arguments, call the repository, `return await reply(run)` . Keep the `-> dict[str, Any]` annotation.
3. Add a Pydantic model in `models.py` if the shape is new.
4. Record a fixture in `tests/fixtures/` and test the repository with a fake source.

Adding a whole data domain

A new domain is **a new repository plus thin tools** — never fetch/parse orchestration in `server.py` . Concretely:

1. Add the narrow Protocol to `sources.py` .
2. Add the fetch method to `client.py` , using the correct dialect (§11).
3. Add the parser to `parsers.py` . Parse by **<thead> label, never column position**.
4. Add the repository, the models, and the tools.
5. Capture a real fixture; tests never touch the network.

Checklist before opening a PR

- `uv run ruff check .` · `uv run mypy src` · `uv run pytest` all clean
- New endpoint → new fixture, and `docs/endpoints.md` updated
- Shape drift raises `EndpointChangedError` rather than returning partial data
- All reads go through `FreezeService.get()` — **never call `PseEdgeClient` from a tool**
- Bump `version` in `pyproject.toml` and roll `CHANGELOG.md` 's Unreleased section **in the same PR** as the work being released

Branching

All work lands on `develop` . `develop` reaches `main` only by pull request; `main` is protected. **Merging to `main` publishes a multi-arch image**, and a merge that changes `version` also cuts a GitHub Release and an immutable `<version>` tag. Never commit to `main` directly.

11. Debugging guide

Symptom → cause

Symptom	Almost certainly
POST /mcp → 421 Invalid Host header	The SDK's DNS-rebinding guard. PSE_PUBLIC_URL must match the host clients actually use — asgi.transport_security_for() derives the allowlist from it
OAuth completes, then the first tool call fails	Same as above. Everything about auth is fine; look at the transport layer
MARKET_OPEN_NO_CACHE	Working as designed. Retry after the close
ENDPOINT_CHANGED	PSE Edge changed shape. Re-capture the fixture, compare against docs/endpoints.md, fix the parser. Never paper over it
HTTP 415 from an .ax endpoint	Wrong dialect — that endpoint needs a JSON body, not form encoding
Disclosure results in the wrong order	sortType must be the literal string "date"
Signup → 503	The mail provider refused. The log names the sender address; ZeptoMail verifies exact domains, so a verified example.com does not cover sub.example.com
Passkeys "just don't work"	PSE_PUBLIC_URL does not match the browser's origin. Enrolled credentials cannot be migrated, only re-enrolled
App container stopped, empty log	A start-time failure, not a crash — a crash always leaves logs. Usually a taken host port, or a bind-mount whose source does not exist
NAS UI shows the project as "Error"	migrate is a one-shot that exits 0. That is success. Judge health by curl /health
ImportError on sqlalchemy in a lean install	Something imported Postgres code eagerly. It must be lazy, inside build_storage()
POST /oauth/token → unauthorized_client	The client is not a machine client. Only pse-edge-admin create-machine-client can authorize this grant — a self-registered client never can, by design
POST /oauth/token → invalid_client (401)	Wrong secret, unknown client, or a revoked one. All answered identically on purpose, so the endpoint is not an oracle for which client ids exist
POST /oauth/token → slow_down (429)	Rate limit: 20/minute per client_id and per IP. Honour Retry-After
A machine client's token stops working early	revoke-machine-client revokes outstanding tokens too, and revocation lands within PSE_TOKEN_CACHE_TTL (60 s)

Useful commands

```
curl -s https://<host>/health # liveness
curl -s https://<host>/health/ready # readiness (checks DB)
curl -s https://<host>/well-known/oauth-protected-resource # is PSE_PUBLIC_URL right?
curl -sS -o /dev/null -w '%{http_code}\n' -X POST https://<host>/mcp # expect 401

docker compose -f compose.nas.yaml logs -f app
docker compose -f compose.nas.yaml exec app pse-edge-admin list-users
docker compose -f compose.nas.yaml exec app pse-edge-admin list-machine-clients
```

Reading the log

Every line is timestamped (ISO-8601 with an offset) in both formats; `PSE_LOG_JSON=1` gives one JSON object per line for a log shipper. Secrets are redacted in both formats as a backstop — nothing here logs a credential deliberately.

The lines worth knowing by sight:

Line	Means
<code>starting pse-edge-mcp <version> ...</code>	the config the process actually resolved to. First thing to read in an incident
<code>upstream: fetching from PSE Edge key=...</code>	a real request left for PSE Edge
<code>upstream: fetched and cached key=... duration_ms=...</code>	it came back, and how slow it was
<code>freeze: refusing an uncached read while the market is open</code>	the policy working, not a fault
<code>auth: rejected a bearer token / quota: refused</code>	refusals only — successes are the access log
<code>oauth: issued a client_credentials access token</code>	a machine client minted a token
<code>oauth: REFRESH TOKEN REUSE detected (WARNING)</code>	a rotated token was presented twice, so it leaked
<code>oauth: DENIED client_credentials to a non-machine client (WARNING)</code>	a misconfigured integration, or someone probing the gate

`upstream: fetching during market hours means the freeze invariant is broken.` Those lines should be a trickle after each 15:00 Manila close and absent between 09:30 and 15:00 on a trading day. That is the single most useful thing to grep this log for.

Two PSE Edge dialects

Getting this wrong produces an HTTP 415 that looks like a server fault:

JSON dialect	POST /common/DisclosureCht.ax Content-Type: application/json → JSON body. Form encoding gives 415.
Form dialect	POST /announcements/search.ax POST /companyDisclosures/search.ax POST /keyword/search.ax Content-Type: application/x-www-form-urlencoded → an HTML FRAGMENT, not JSON.
Wire dates are MM-dd-yyyy. PSE Edge's parameter names sometimes differ from its own HTML form field names – trust docs/endpoints.md over the page source.	

12. Testing

Tests never touch PSE Edge. All HTTP is mocked with `respx` against 15 fixtures in `tests/fixtures/`, recorded from real captures. A new endpoint requires a new fixture.

Layer	How it is tested
Repositories	Five-line fake sources — no HTTP mocking (<code>test_repositories.py</code>)
Client / parsers	<code>respx</code> + recorded fixtures
Freeze policy	Injected calendar pinning the clock (<code>test_service.py</code>)
Postgres	testcontainers, marked <code>@pytest.mark.postgres</code>
Auth	<code>test_auth_journey.py</code> — soft-webauthn, real signatures , real Postgres
Deployment	<code>test_deploy.py</code> — health probes, logging, the real ASGI factory

A lesson worth internalizing

250 passing tests once missed a bug that made the server unusable in production: the transport-security setting. The journey test passed `transport_security=...` **explicitly**, so it exercised a configuration production never built — the fixture worked around the exact defect it should have caught.

A fixture that configures something production leaves at its default is a blind spot, not a test. Where practical, test through `create_app()` — the real factory — rather than hand-assembling the stack.

13. Configuration reference

All settings live in `config.py` as a frozen dataclass; environment variables override.

Variable	Default	Notes
PSE_EDGE_BASE_URL	https://edge.pse.com.ph	
PSE_MARKET_OPEN / PSE_MARKET_CLOSE	09:30 / 15:00	Asia/Manila; drives the freeze
PSE_THROTTLE_RPS / PSE_THROTTLE_BURST	1.0 / 2	Outbound politeness
PSE_TIMEOUT_SEC / PSE_RETRY_ATTEMPTS	20.0 / 3	
DATABASE_URL	unset	Unset ⇒ in-memory + no archive
PSE_DB_POOL_SIZE / PSE_DB_MAX_OVERFLOW	5 / 10	
PSE_AUTH_REQUIRED	false	Requires DATABASE_URL
PSE_TOKEN_CACHE_TTL	60	The revocation-latency budget
PSE_QUOTA_PER_MIN / PSE_QUOTA_PER_DAY	60 / 2000	Per user, in-process
PSE_PUBLIC_URL	http://localhost:8000	The highest-risk setting — see below
PSE_ACCESS_TTL_MIN / PSE_REFRESH_TTL_DAYS	30 / 30	
ZEPTOMAIL_API_KEY	unset	Unset ⇒ emails logged to console
PSE_EMAIL_FROM	no-reply@localhost	Must be a verified sender domain
PSE_USAGE_RETENTION_DAYS	90	The privacy page states 90
PSE_STATEFUL / PSE_SSE	false / false	See below
PSE_LOG_JSON / PSE_LOG_LEVEL	false / INFO	

PSE_PUBLIC_URL must be the real external HTTPS URL. It simultaneously drives the WebAuthn `rp_id`, the links in verification emails, the OAuth issuer in discovery documents, and the DNS-rebinding host allowlist. A wrong value breaks passkeys in a way that looks like a browser bug, and enrolled credentials cannot be recovered — only re-enrolled. Verify it after every deploy: `curl -s https://<host>/.well-known/oauth-protected-resource`

Why HTTP is stateless with plain JSON by default

The server is read-only tools over data the freeze holds still. It uses **none** of the features MCP sessions exist to enable — no notifications, no resource subscriptions, no sampling, no elicitation, no progress. Every request is self-contained.

That makes horizontal scaling ordinary: any replica serves any request behind plain round-robin — no sticky routing, no per-session memory, no event store. Without SSE, idle clients hold no connection, so N users stop meaning N concurrent connections.

`--stateful` and `--sse` restore session mode for anyone who needs resumability or server-initiated messages. They are independent flags: statelessness is the scaling property, plain JSON is what keeps ordinary proxies out of the way.

Workers are processes. All in-memory state — quota windows, the parse memo, the token cache, the usage buffer — is **per worker**, so N workers means a user's effective quota ceiling is up to N× nominal. Scale workers for CPU, and scale limits with them.

14. Deployment topologies

A. VPS that owns ports 80/443

compose.prod.yaml

internet
| :80 / :443
▼




Requires the router to forward
80→8280 and 443→8243: ACME
validates on 80/443 only.

B. NAS behind Cloudflare Tunnel

compose.nas.yaml [+ compose.tunnel.yaml]

internet
|
▼ (Cloudflare edge, TLS ends here)



Needs no inbound ports at all.
Tunnel route is HTTP → app:8000
(container port, not the host port)

Both pull the published, CI-gated multi-arch image; neither builds from source. Pin `PSE_IMAGE_TAG` — a moving `:latest` that lags the compose file reads as a broken deployment.

The runtime image contains only what the app needs to run: multi-stage, non-root, no build tools, no dev dependencies, no source tree, no credentials in any layer. `scripts/check_image.py` asserts installed distributions **equal** the resolved runtime closure from `uv export`, so any stray package fails the build. Size is reported as information only — **the gate is necessity, not a number**.

15. Gotchas

Findings that cost real debugging time. Most are recorded in `docs/endpoints.md` with the evidence.

- **Parse disclosure tables by `<thead>` label, never by column position.** The two search endpoints order their columns differently.
- **Financial-report units labels contradict each other** between the annual and quarterly sections. Values are passed through **verbatim and never rescaled**; each period reports its own `currency_units` for the caller to check.
- **Index changes are printed unsigned.** PSE Edge shows direction only as a colour and a ▲/▼ glyph, so signs are derived from the glyph.
- **selectolax traps:** `iter()` walks direct children only, and `css("a, b")` returns matches **grouped by selector, not in document order** — which once filed all four financial statements under the last heading. Use `root.traverse()`.
- **Accounting negatives** `(1,234)` parse to `-1234`.
- **dividends_and_rights_form.do is an empty shell.** It posts to `dividends_and_rights_list.ax` once per tab, with `DividendsOrRights` in the query string and `cmpy_id` in the body.
- **Homepage feeds are keyed by PSE Edge's own group labels**, not invented buckets.
- **market_calendar.PSE_HOLIDAYS is a yearly-maintained seed.** Verify it against the PSE holiday circular each December and whenever touching calendar logic.

- `/health` must never touch the database. A DB-dependent liveness probe restarts every replica during a blip, turning a recoverable outage into an outage plus a restart storm. `/health/ready` is where the DB check belongs.

16. Where to look next

I want to...	Read
Understand why a decision was made	<code>docs/plan.md</code>
Add or fix an upstream call	<code>docs/endpoints.md</code> , then <code>client.py</code> + <code>parsers.py</code>
Change what a tool returns	<code>models.py</code> , then the repository
Change caching behaviour	<code>service.py</code> — and re-read §4 first
Deploy or operate it	<code>docs/deploy.md</code>
Know the invariants quickly	<code>CLAUDE.md</code>